

Lab 1: Descriptive Statistics

due date: Wednesday, September 6 in class

Introduction

Below are some data on violent crime in US States during the year 1973. This data set comes from R, a free statistics software package widely used by statisticians.

You are asked to compute some summary statistics, create boxplots and histograms, and answer a few questions. All of this work should be done with R or RStudio. (Part of your grade is simply to complete this lab using the software package R.) Information about the website where you can download R and some basic R commands have been included at the end of this handout to help. After starting RStudio and loading the data set `USArrests`, type `help(USArrests)` to get some information about this dataset.

For submission to your instructor, write your answers in the space provided and hand in graphs.

Data

Violent Crime Rates by US State

Description:

This data set contains statistics, in arrests per 100,000 residents for assault, murder, and rape in each of the 50 US states in 1973. Also given is the percent of the population living in urban areas.

Questions

1. What is the population under study?
2. Compute the Min, Q1, Median, Q3, Max for the numerical variables Murder and Rape. Then compute the mean for these two variables. Is the mean bigger than, smaller than, or roughly equal to the median? What does your answer here tell you about the variables Murder and Rape.
3. Compute boxplots for the variables Murder and Rape. Summarize what you learn from viewing the boxplots for the two variables. Plots of these boxplots should be handed in with your lab.

4. Looking at the data for Alaska, how does it compare to other states? Do you think Alaska is typical? Can you think of any factors that might explain the data for Alaska in comparison with the other states?

5. Make a frequency histogram for the numerical variable Rape. Let each bin be of size 5, and have the range on the x -axis be from 0 to 50. (This means your bins will be $[0, 5]$, $[5, 10]$, etc.) How many states have between 15 and 20 rapes per 100,000 residents per year? Write the R command you used to plot this histogram in the space below.

6. Looking at the histogram for the variable Rape, can you tell if the mean is bigger than the median? Explain.

7. Now load the data table in the file `lab1_examScores.dat` into R and compute the mean and median, and plot a histogram. This data has been sorted for your convenience.
Remove the largest test score and save it to a new variable called `d`. Type

```
d = testScores[1:99]
```


at the R prompt to do this. Explain the effect of removing the largest test score on the mean and the median.

R and RStudio are available over the internet. The URL for the Comprehensive R Archive Network is <http://www.r-project.org/>

but a simple Google search will find you plenty of hits. You will need to choose a mirror and download a package appropriate for your operating system. If you are not familiar with downloading such a software package, see me during office hours for help.

All commands should be typed in R's console window.

Helpful commands from R:

> help(?????)	get help on command ?????
> data(USArrests)	loads data set USArrests
> USArrests	displays dataset
> attach(USArrests)	makes variables Murder, Assault, UrbanPop, Rape available by name
> help(USArrests)	get info on dataset
> Murder	display variable Murder
> summary(??)	display summary statistics for variable ??
> fivenum(??)	display Tukey's five number summary for variable ??
> boxplot(Murder,Rape)	creates boxplots for variables Murder, Rape
> boxplot(USArrest)	creates boxplots for all variables in dataset
> hist(Rape,seq(0,50,5))	creates a frequency histogram for Rape, with bins of size 5 and a range from 0 to 50
>testScores=read.table("lab1_examScores.dat")	loads exam score data
>testScores	displays data
>testScores=testScores[,1])	gets rid of row names